

The Relational Model ER to Relational / SQL

Announcements

Project I Part I: Starting tomorrow!
Partners: Try Piazza

Review

Relation: Set of tuples with typed values (table)
Schema: Names and types for values in relation
Database: Set of relations
SQL: Structured Query Language
Integrity constraint: Restrictions on valid data
Candidate key: Minimal fields that identify tuple
Primary key: Designated identifier for a tuple
Foreign key: Reference to another tuple

Referential Integrity

A database instance has *referential integrity* if all foreign key constraints are enforced *no dangling references*

Examples where referential integrity is not enforced
Web links
Restaurant menus
Some relational databases!

How to Enforce Integrity Constraints

Run checks any time database changes

On INSERT

what if new Enrolled tuple refers to non-existent student?
Reject insertion

On DELETE (many options)

what if Students tuple is deleted?
delete dependent Enrolled tuples
reject deletion
set Enrolled.sid to default value or *null*
(null in SQL: unknown/inapplicable)

How to Enforce Integrity Constraints

Run checks any time database changes

On INSERT

what if new Enrolled tuple refers to non-existent student? *Reject insertion*

On DELETE (many options)

what if Students tuple is deleted?
delete dependent Enrolled tuples (*CASCADE*)
reject deletion (default: *ON DELETE NOACTION*)
set Enrolled.sid to default value or *null*
(null in SQL: unknown/inapplicable)

General Constraints

Boolean constraint expression added to schema
Very powerful; not common

```
CREATE TABLE Enrolled(
  sid int,
  cid int,
  grade char(2),
  CHECK (
    boolean expression
  )
)
```


General Constraints

Boolean constraint expression added to schema
Very powerful; not common

```
CREATE TABLE Enrolled(
  sid int,
  cid int,
  grade char(2),
  CHECK (
    grade = 'A' or grade = 'B' or
    grade = 'C' or grade = 'D' or
    grade = 'F'
  )
)
```

Where do ICs come from?


Based on application semantics and use cases
IC is statement about the world (all possible instances)
Can't infer ICs by staring at an instance
e.g. Is "login" a unique candidate key? Maybe?
Unique and foreign key ICs are very common
(others less so)



Translating ER → Relational Models

Entity Set → Relation


Include all attributes
Entity set key become relation primary key



```
CREATE TABLE Course(
  cid int,
  name text,
  loc text,
  PRIMARY KEY (cid)
)
```

Relationship Set w/out constraint → Relation

Add keys for each entity set as foreign keys: *superkey*
(if there are other constraints: may not need all columns)
All attributes of the relationship set




```
CREATE TABLE Takes(
  uid int,
  cid int,
  since date,
  PRIMARY KEY (uid, cid),
  FOREIGN KEY (uid) REFERENCES Users,
  FOREIGN KEY (cid) REFERENCES Courses
)
```

At Most One → Relation

Add relationship attributes (none here)

Add keys for entity set as foreign keys

What is the primary key?




At Most One: Combine?

Zero or 1 courses, cid is primary key; Similar to ???

Combine Instructs attributes into Courses (preferred)

How to represent courses without instructor?

NULL uid (and other Instructs attributes)




Exactly One Constraint → Relation

Represent: Course must have Instructor?

Combine relationship into Courses + NOT NULL


What happens if we delete User who is Instructor?

Default ON DELETE NO ACTION: Don't permit



At Least One Constraint?

```
CREATE TABLE Instructs(
  cid int,
  uid int,
  PRIMARY KEY (cid, uid),
  FOREIGN KEY (cid) REFERENCES Course,
  FOREIGN KEY (uid) REFERENCES User)
```



Can't be easily expressed!

Weak Entity → Relation


Weak entity set and identifying relationship set are translated into a single table.

When the owner entity is deleted, all owned weak entities must also be deleted.



ISA Hierarchies

```
CREATE TABLE Users(uid int, name text, PRIMARY KEY(uid))
CREATE TABLE Instructors(uid int, rating int,
  PRIMARY KEY(uid), FOREIGN KEY (uid) REFERENCES Users)
CREATE TABLE Students(uid int, grade char(2),
  PRIMARY KEY(uid), FOREIGN KEY (uid) REFERENCES Users)
```



ISA Hierarchies

Option 1: Keep base relation

Instructors & Students recorded in Users


Extra info in Instructors or Students relation

Primary key = key of parent table; is foreign key for child tables


Option 2: Only keep child relations

Children copy attributes from Users

Only if covering
constraint = yes
(all users are
instructors or
students)



Aggregation




Aggregation

Convert the aggregated relationship into an entity

Convert like any other entity

E.g. Donates: PRIMARY KEY (Company_id, Course_id)

Manages: References this key



ER-Relational Review

Not complete. Just tips

Postgres documentation

Good reference on SQL

<https://www.postgresql.org/docs/9.3/static/ddl.html>

The textbook is actually quite good on this topic!

Feedback

<http://tinyurl.com/w4lll>