

Administrative Stuff

Lots of FD exercises.

AA3 Due this morning. How was it?

AA4 released, due in 1 week. Window functions

Reminder: my OH moved to Tues after class

L8

Disk, Storage, and Indexing


Eugene Wu

Fall 2018

DBMS Overview

Each layer provides a simple abstraction to layers above it, and makes assumptions about layers below it.


Requires careful design and assumptions for performance



DBMS Overview

Applications interact with
SQL Client

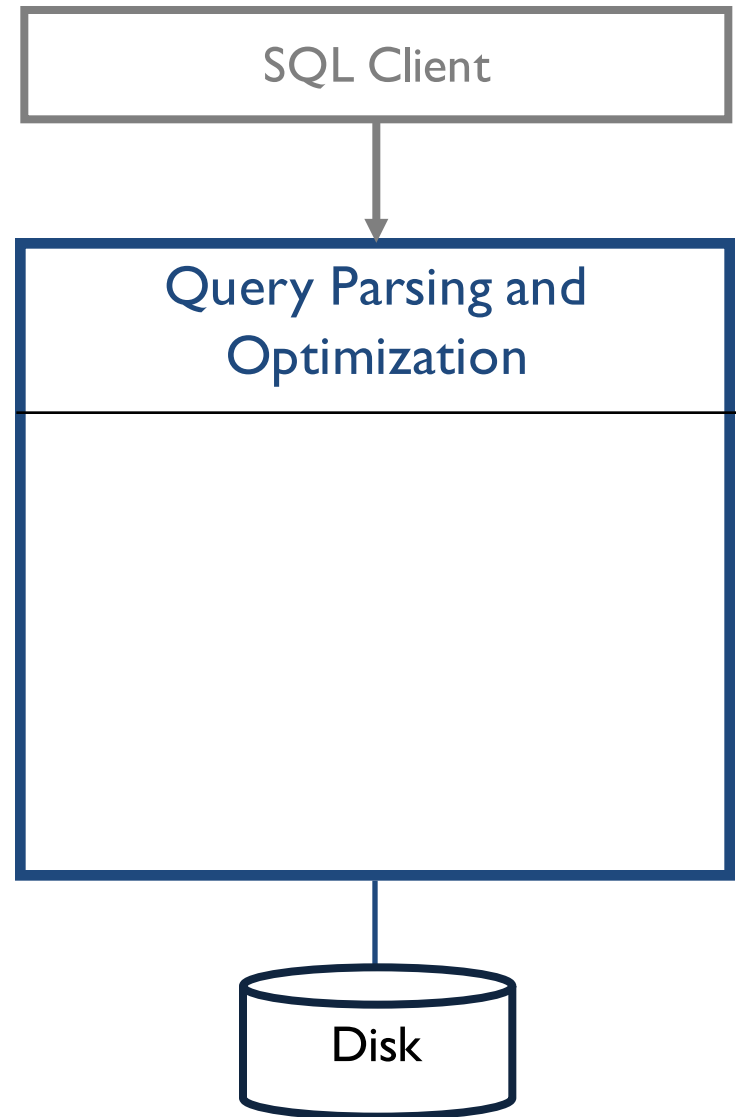

```
db.execute('''  
    SELECT a, b  
    FROM S, T  
    WHERE S.c = T.c''')
```



DBMS Overview

Parse, check, and verify the SQL query

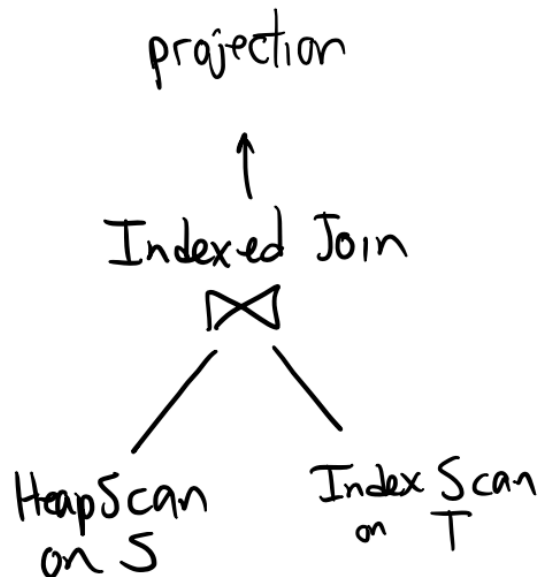
Turn into efficient query plan



DBMS Overview

A query is represented as a relational data flow

Each operator is a specific implementation



DBMS Overview

Organizes tables, indexes, records as groups of pages in a “logical file”

API:

- Operators ask for records
- Logical files help read and write bytes on pages



DBMS Overview

Not all pages can fit into RAM.

Buffer manager provides illusion that all pages are accessible.

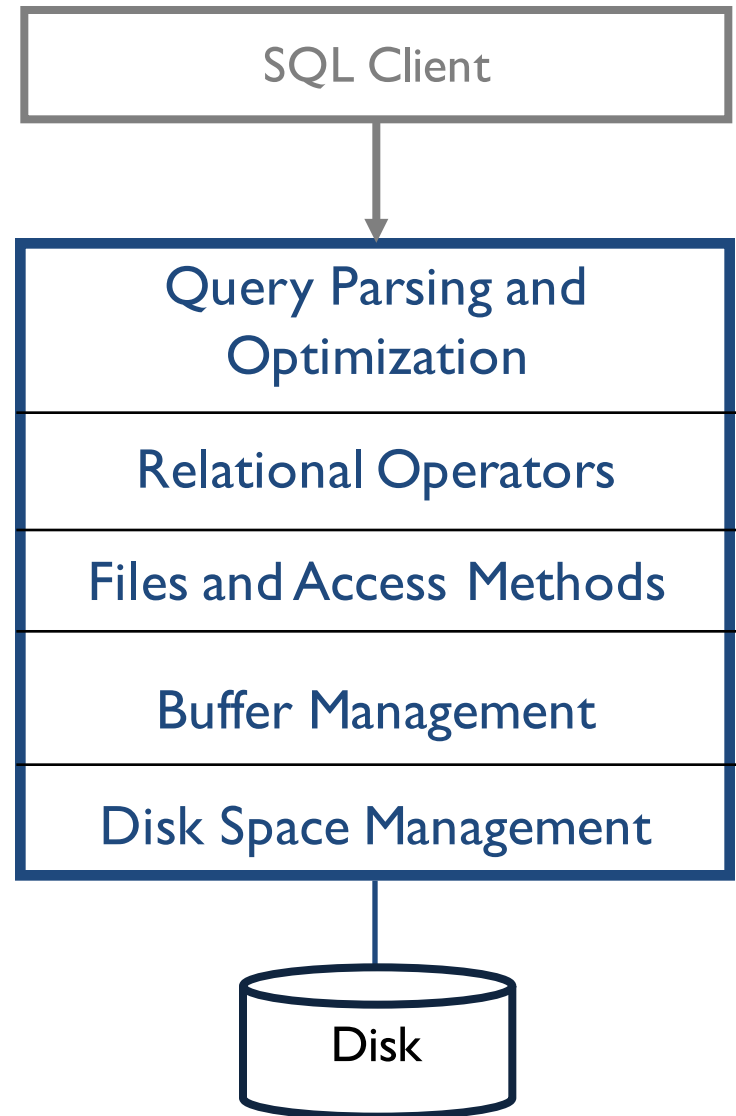
Files simply ask for pages.



DBMS Overview

Physically read and write bytes on one or more storage devices (hard drives, SSDs, etc)

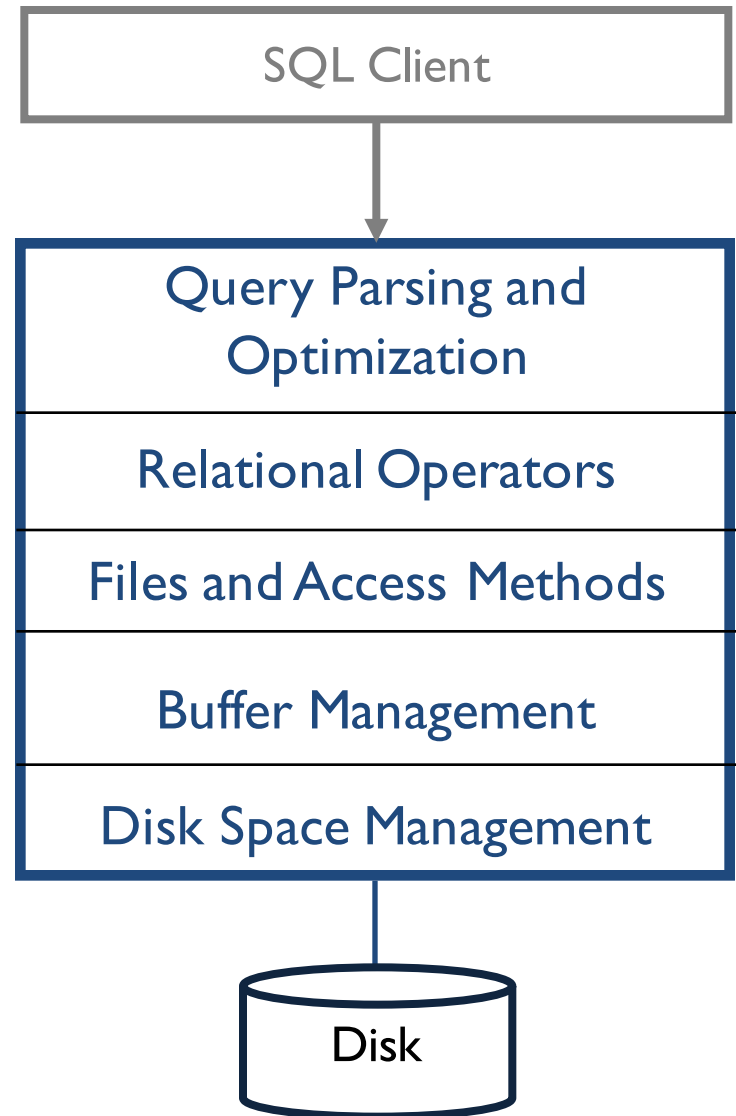
Storage performance properties dictate the design of layers above.



DBMS Overview

Layers help manage engineering complexity.


Requires assumptions about performance of lower layers. (cost models)



DBMS Overview

Physically read and write bytes on one or more storage devices (hard drives, SSDs, etc)

Storage performance properties dictate the design of layers above.



Storage Devices

Hard Drives vs SSDs vs RAM

Locality (random vs sequential)

Reads vs Writes

Performance and Monetary Costs

\$ Matters

Why not store all in RAM?

Costs too much

High-end Databases today ~Petabyte (1000TB) range.

SQL Hyperscale scales to 100TB+ (announced Sep 24, 2018)

~60% cost of a production system is in the disks.

Main memory not persistent

Obviously important if DB stops/crashes

Some systems are *main-memory* DBMSes, topic for advanced DB course


\$ Matters


Newegg enterprise \$1000

RAM: 0.08TB

SSD: ~2TB (25x)

Disk: ~40TB (500x)







Interesting numbers

compress 1k bytes: 3000 ns



roundtrip in data center: .5 ms



Spin speed: ~7200 RPM

Section = multiples of pages

Terminology: block same as page in this class



Time to access (read or write) a disk block

seek time 2-4 msec avg

rotational delay 2-4 msec

transfer time 0.3 msec/64kb page

Throughput


read ~150 MB/sec

write ~50 MB/sec

Key: reduce seek and rotational delays

HW & SW approaches

of 4 byte values read per second



5 orders of magnitude

SSDs


NAND memory

Small reads: 4-8k

Big writes: 1-2MB

2-3k writes before failure

wear leveling: distribute writes around



Write amplification: changing 4 bytes writes a 1-2MB chunk!

Need to think about wear, garbage collection, writing in bulk


SSDs

Reads fast

- single read time: 0.03ms
- random reads: 500MB/s
- sequential reads: 525MB/s

Writes less predictable

- single write time: 0.03ms
- random reads: 120MB/s
- sequential reads: 480MB/s



What's Best? Depends on Application

Small databases:

- All global daily weather since 1929: 20GB

- 2000 US Census: 200GB

- 2009 english wikipedia: 14GB

- Easily fits on an SSD or in RAM

Very Big databases


- Sensors easily generate TBs of data/day

- Boeing 787 generates 1/2 TB per flight

- Disk has best cost-capacity ratio

- SSDs help reduce read variance

Work from the bottom up



Strategies for Fast Data Access

Big difference between random & sequential access

- Optimize for sequential accesses

Amortize sequentially read & write big chunks of bytes

Cache popular blocks

Pre-fetch what you will need later

API

- read/write page
- read/write sequential pages
- notion of “next” page (upper layers can assume next is faster)

Big File

Files and Access Methods


Disk Space Management

OS File System

Disk

DBs don't directly talk to storage devices.

Usually go through OS file system



Big File

Files and Access Methods

Disk Space Management

Big File
part 1

OS File System

Disk

Big File
part 2


OS File System

Disk

Big File
part 3

OS File System

Disk




PostgreSQL database stored in files

```
ewu@dyn-160-39-248-197 ~/d/p/base> cd 33883/
ewu@dyn-160-39-248-197 ~/d/p/b/33883> ls
112      12526      2601      2612      2659      2696      2840_vm   3456_fsm   3602
113      12528      2601_fsm  2612_fsm  2660      2699      2841      3456_vm   3602_fsm
1247     12529      2601_vm   2612_vm   2661      2701      2995      3466      3602_vm
1247_fsm 12529_fsm  2602      2613      2662      2702      2995_vm   3466_vm   3603
1247_vm   12529_vm  2602_fsm  2613_vm   2663      2703      2996      3467      3603_fsm
1249     12531      2602_vm   2615      2664      2704      3079      3468      3603_vm
1249_fsm 12533      2603      2615_fsm  2665      2753      3079_fsm  3501      3604
1249_vm   12534      2603_fsm  2615_vm   2666      2753_fsm  3079_vm   3501_vm   3605
12504     12536      2603_vm   2616      2667      2753_vm   3080      3502      3606
12504_fsm 12538      2604      2616_fsm  2668      2754      3081      3503      3607
12504_vm   1255      2604_vm   2616_vm   2669      2755      3085      3534      3608
12506     1255_fsm  2605      2617      2670      2756      3118      3541      3609
12508     1255_vm   2605_fsm  2617_fsm  2673      2757      3118_vm   3541_fsm  3712
12509     1259      2605_vm   2617_vm   2674      2830      3119      3541_vm   3764
12509_fsm 1259_fsm  2606      2618      2675      2830_vm   3164      3542      3764_fsm
12509_vm 1259_vm   2606_fsm  2618_fsm  2678      2831      3256      3574      3764_vm
12511     1417      2606_vm   2618_vm   2679      2832      3256_vm   3575      3766
12513     1417_vm   2607      2619      2680      2832_vm   3257      3576      3767
12514     1418      2607_fsm  2619_fsm  2681      2833      3258      3576_vm   548
12514_fsm 1418_vm   2607_vm   2619_vm   2682      2834      3394      3596      549
12514_vm   174      2608      2620      2683      2834_vm   3394_fsm  3596_vm   826
12516     175      2608_fsm  2620_vm   2684      2835      3394_vm   3597      826_vm
12518     2187      2608_vm   2650      2685      2836      3395      3598      827
12519     2328      2609      2651      2686      2836_vm   34002     3598_vm   828
12519_fsm 2328_vm   2609_fsm  2652      2687      2837      34004     3599      PG_VERSION
12519_vm 2336      2609_vm   2653      2688      2838      34004_fsm 3600      pg_filenode.map
12521     2336_vm   2610      2654      2689      2838_fsm  34004_vm  3600_fsm  pg_internal.init
12523     2337      2610_fsm  2655      2690      2838_vm   34008     3600_vm
12524     2600      2610_vm   2656      2691      2839      34031     3601
12524_fsm 2600_fsm  2611      2657      2692      2840      3455      3601_fsm
12524_vm 2600_vm   2611_vm   2658      2693      2840_fsm  3456      3601_vm
ewu@dyn-160-39-248-197 ~/d/p/b/33883>
```

Files

Logical relation



Files

Higher layers want to talk in terms of records, and files of records

File: collection of pages

Minimum API:

- insert/delete/modify record

- lookup record_id

- scan all records

Page: collection of records

- typically *fixed page sizes* (8 or 64kb in PostgreSQL)

These are logical.

Different page organizations in a file have different access costs

Units that we'll care about

Ignore CPU cost

Ignore RAM cost

B # *data* pages on disk for relation

R # records per data page

D avg time to read/write data page to/from disk

Simplifies life when computing costs

OK to not be exactly correct

Unordered Heap Files

Collection of records (no order)

As we add records, pages allocated

As we remove records, pages removed

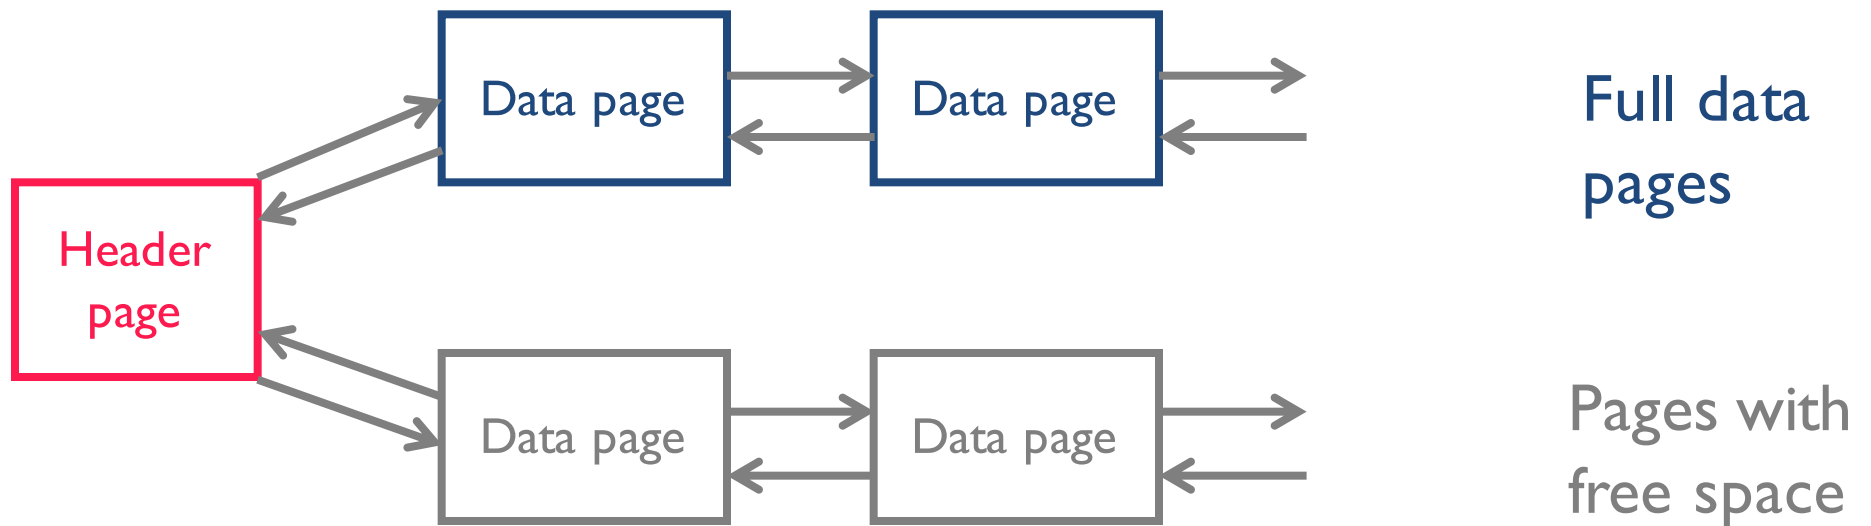
To support record level ops, need to track:

- pages in file

- free space on pages

- records on page

Heap File




Header page info stored in catalog

Data page contains: 2 pointers, free space, data

Need to scan pages to answer any query

Which page has enough free space for 100 bytes?

Use a directory



Directory entries track # free bytes on data pages

Directory is collection of pages

Indexes

“If I had eight hours to chop down a tree,
I'd spend six sharpening my ax.”

Abraham Lincoln

Indexes

Heap files answers any query via a sequential scan

Queries use *qualifications* (predicates)

find students here class = “CS”

find students with age > 10

Indexes: file structures for value-based queries

B+-tree index (~1970s)

Hash index

Overview! Details in 4112

Indexes

Defined wrt a *search key*


no relation to candidate keys!

Faster access for WHERE clauses w/ search key

```
CREATE INDEX idx1 ON users USING btree (sid)
CREATE INDEX idx2 ON users USING hash  (sid)
CREATE INDEX idx3 ON users USING btree (age,name)
```

You will play around with indexes in HW4

Primary vs Secondary Index Files



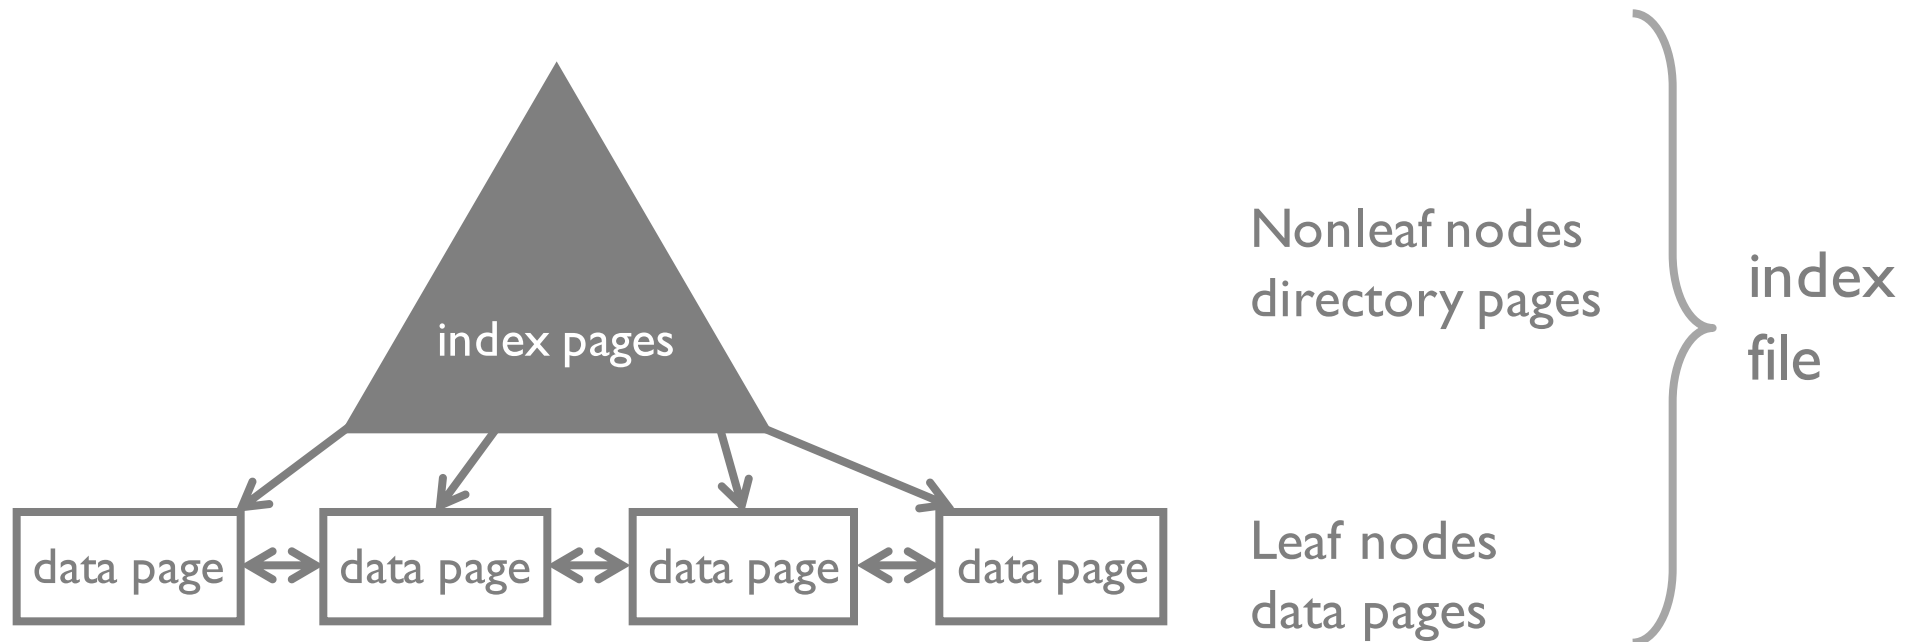
Primary: data entries contain data records in another file

Pros: directly access data

Secondary: data entries contain pointers $\langle \text{search key value, rid} \rangle$

Pros: index is more compact

B+ Tree Index



Supports equality and range queries on search key

Self balancing

Leaf nodes are connected

Disk optimized

Node = Page

Height is with respect to directory pages (the gray part of the triangle)

B+ Tree on (age)

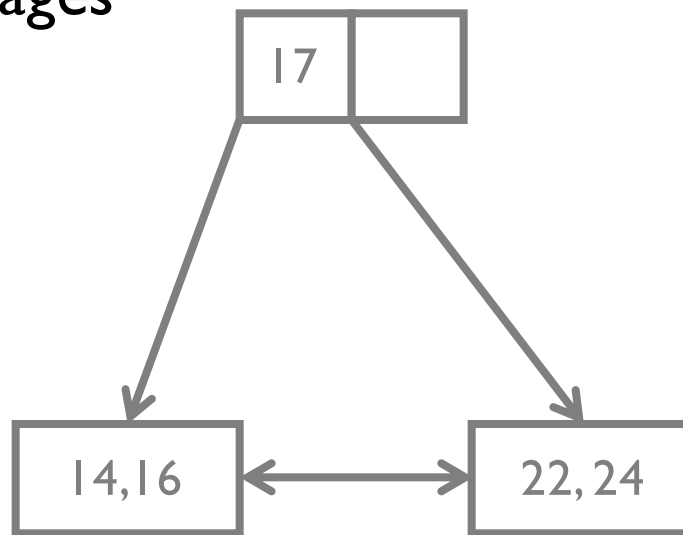
Non-leaf Directory pages

m index entries

m+1 pointers

Leaf Data pages

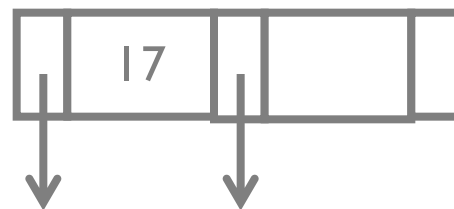
data entries/tuples



index & data
page contents
are in order

Query: `SELECT * WHERE age= 14`

directory page



Index Only Queries: B+ Tree on (age)


Non-leaf Directory pages

m index entries

m+1 pointers

Leaf Data pages


data entries/tuples



index & data
page contents
are in order


Query: **SELECT age WHERE age > 14**
(index only!)

directory page




B+ Tree on (age)

Note: 50 not a
data entry




Query: `SELECT * WHERE age = 55`


directory page




B+ Tree on (age)



B+ Tree on (age)




B+ Tree on (age)



Query: **SELECT * WHERE age > 20**

B+ Tree on (age, name)



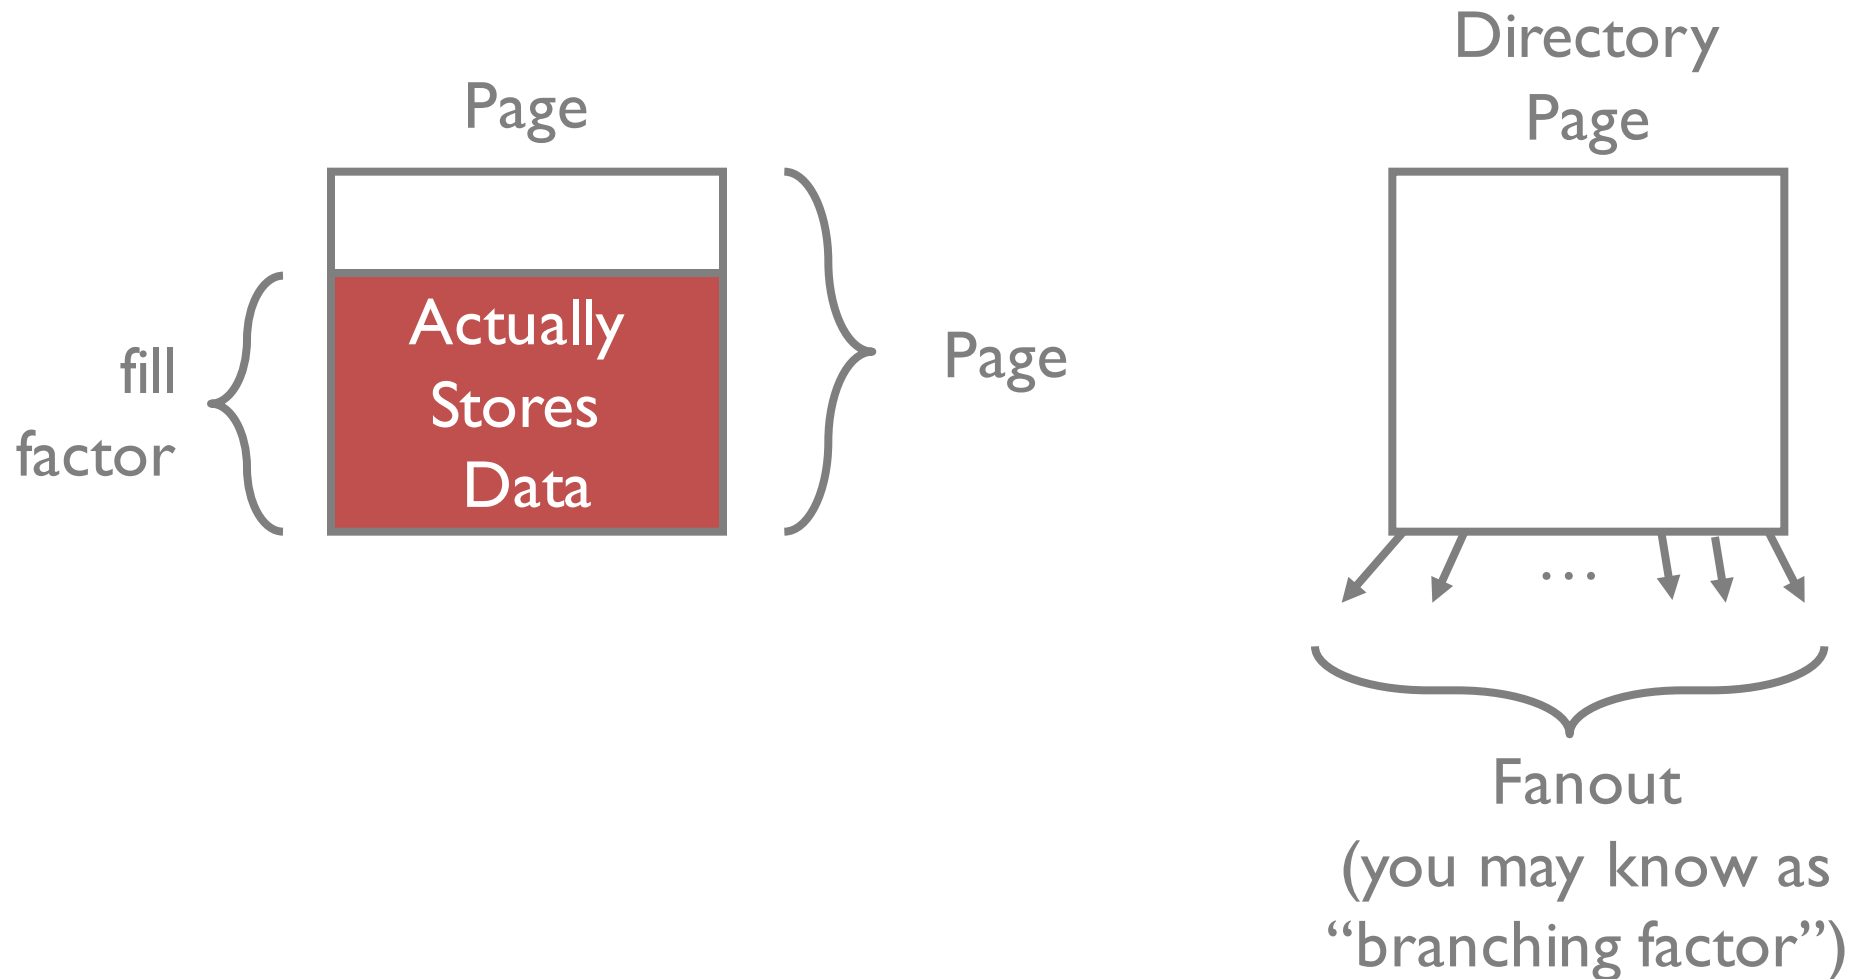
How do the following queries use the index on (age, name)?

SELECT age WHERE age = 14

SELECT * WHERE age < 18 AND name < 'monica'

SELECT age WHERE name = 'bobby'

Terminology



Some numbers (8kb pages)

How many levels?

fill-factor: ~66%

~300 entries per directory page

height 2: $300^3 \sim 27$ Million entries

height 3: $300^4 \sim 8.1$ Billion entries

Top levels often in memory

height 2 only 300 pages ~2.4MB

height 3 only 90k pages ~750MB

Hash Index on age

Hash function

$$h(v) = v \% 3$$


Hash buckets
containing
data pages

0, 6

4, 13

5, 14

INSERT Hash Index on age



INSERT Hash Index on age

Search key

11

Hash function

$$h(v) = v \% 3$$

Hash buckets
containing
data pages

0, 6


4, 13

5, 14




1

INSERT Hash Index on age



SEARCH Hash Index on age



Good for equality selections

Index = data pages + overflow data pages

Hash function $h(v)$ takes as input the *search key*

Costs

Three file types

Heap, B+ Tree, Hash

Operations we care about

Scan all data `SELECT * FROM R`

Equality `SELECT * FROM R WHERE x = I`

Range `SELECT * FROM R WHERE x > 10 and x < 50`

Insert record

Delete record

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything				
Equality				
Range				
Insert				
Delete				

B # data pages

D time to read/write page

M # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD			
Equality	0.5BD			
Range	BD			
Insert	2D			
Delete	Search + D			

Heap File

equality on a key. How many results?

B # data pages

D time to read/write page

M # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD		
Equality	0.5BD	$D(\log_2 B)$		
Range	BD	$D(\log_2 B + M)$		
Insert	2D	Search + BD		
Delete	Search + D	Search + BD		

Heap File

equality on a key. How many results?

Sorted File

files compacted after deletion

B # data pages

D time to read/write page

M # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD	1.25BD	
Equality	0.5BD	$D(\log_2 B)$	$D(\log_{80} B + 1)$	
Range	BD	$D(\log_2 B + M)$	$D(\log_{80} B + M)$	
Insert	2D	Search + BD	$D(\log_{80} B)$	
Delete	Search + D	Search + BD	$D(\log_{80} B)$	

Heap File

equality on a key. How many results?

Sorted File

files compacted after deletion

B+ Tree

100 entries/directory page

80% fill factor

B # data pages

D time to read/write page

M # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD	1.25BD	1.25BD
Equality	0.5BD	$D(\log_2 B)$	$D(\log_{80} B + 1)$	D
Range	BD	$D(\log_2 B + M)$	$D(\log_{80} B + M)$	1.25BD
Insert	2D	Search + BD	$D(\log_{80} B)$	2D
Delete	Search + D	Search + BD	$D(\log_{80} B)$	2D

Heap File

equality on a key. How many results?

Sorted File

files compacted after deletion

B+ Tree

100 entries/directory page

80% fill factor

Hash index

no overflow

80% fill factor

B # data pages

D time to read/write page

M # pages in range query

Where do B, D, M come from?

Estimated from more basic info

Assuming:

- fanout = number of directory entries
- pointer in secondary index same size as directory entry

Given:

- p: page size
- r: record size
- d: directory entry size
- f: fill factor
- n: # records

Estimate for primary and secondary index:

- size
- height
- access cost

Primary B+ Index

$p = 100$	records/page	$p / r = 10$
$r = 10$	direntries/page	$p / d = 20$
$d = 5$	fanout:	20
$f = 100\%$	# data pages	$n/(p/r) = 800$
$n = 8000$	height	$\log_{20} 1000 = 3$

Cost to look up a single record is
3 for directory pages + 1 data page

Secondary B+ Index

$p = 100$	records/page	$p / r = 10$
$r = 10$	direntries/page	$p / d = 20$
$d = 5$	fanout:	20
$f = 100\%$	# data pages	$n/(p/d) = 400$
$n = 8000$	height	2

Cost to look up a single record is
1 for directory pages + 1 data page + 1 pointer lookup

How to pick?

Depends on your queries (workload)

Which relations?

Which attributes?

Which types of predicates ($=$, $<$, $>$)

Selectivity

Insert/delete/update queries? how many?

How to choose indexes?

Considerations

which relations should have indexes?

on what attributes?

how many indexes?

what type of index (hash/tree)?

Naïve Algorithm

get query workload

group queries by type

for each query type in order of importance

 calculate best cost using current indexes

 if new index IDX will further reduce cost

 create IDX

Why not create every index?

update queries slowed down (upkeep costs)

takes up space

High level guidelines

Check the WHERE clauses

- attributes in WHERE are search/index keys

- equality predicate → hash index

- range predicate → tree index

Multi-attribute search keys supported

- order of attributes matters for range queries

- may enable queries that don't look at data pages (*index-only*)

Summary

Design depends on economics, access cost ratios

Disk still dominant wrt cost/capacity ratio

Many physical layouts for files

- same APIs, difference performance

- remember physical independence

Indexes

- Structures to speed up read queries

- Multiple indexes possible

- Decision depends on workload

Things to Know

- How a hard drive works and its major performance characteristics
- The storage hierarchy and rough performance differences between RAM, SSD, Hard drives
- What files, pages, and records are, and how they are different than the UNIX model
- Heap File data structure
- B+ tree and Hash indexes
- Performance characteristics of different file organizations